

R Notebook

Projet Unsupervised Learning

Let's Simulate a directed graph with $K = 8$ nodes using Erdos-Renyi's model with $p = 0.4$. We will do so using the package **igraph**: The model can be written probabilistically as follows:

$$X_{ij} | \{Z_i = k, Z_j = l\} \sim \text{Bern}(0.4)$$

where a 1 in the matrix corresponds to website i pointing toward website j . Note that in this setting we will assume that a website cannot point to itself (i.e the diagonal of the matrix X is 0).

```
set.seed(20222023)
library(igraph)
```

```
##
## Attachement du package : 'igraph'


## Les objets suivants sont masqués depuis 'package:stats':
##
##      decompose, spectrum

## L'objet suivant est masqué depuis 'package:base':
##
##      union
```

```
library(ggraph)
```

```
## Le chargement a nécessité le package : ggplot2
```

```
library(ggplot2)
G1 = sample_gnp(n = 8, p = .4, directed = T, loops = TRUE)
X = as.matrix(as_adjacency_matrix(G1))
#pdf('mtg.pdf')
p1=plot(G1)
```



```
#dev.off()
X
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8]
## [1,]    0    1    0    1    0    0    0    1
## [2,]    1    0    1    0    1    1    0    0
## [3,]    0    0    0    0    0    0    1    0
## [4,]    0    1    1    1    0    0    0    0
## [5,]    0    1    0    0    0    0    1    1
## [6,]    1    0    0    1    1    1    0    0
## [7,]    1    0    0    1    0    1    0    0
## [8,]    1    0    0    0    1    1    0    0
```

Let's write the transition matrix using the random surfer's model with probability $p = 1 - \epsilon$ of choosing the available hyperlinks and probability ϵ of choosing a website randomly out of the 8 available websites:

Using the above hypothesis the transition matrix can be written as

$$\mathbf{A} = (1 - \epsilon)\mathbf{M} + \frac{\epsilon}{K}\mathbf{I}_8.$$

Where

$$M_{ij} = \frac{X_{ij}}{\sum_j X_{ij}}, \forall 1 \leq i, j \leq 8;$$

```
TransMat = function(adjMat, epsilon){
  K = ncol(adjMat)
  A = (1- epsilon)* (adjMat / rowSums(adjMat)) + (epsilon/K) * diag(K)
  return(A)
}
```

Let's fix $\epsilon = 0.05$, and assuming that the Markov Chain is aperiodic and irreducible let's compute the stationary distribution:

Recall for an aperiodic and irreducible \mathcal{MC} finding the stationary distribution amounts to solving the equation $\pi \mathbf{A} = \pi$. Furthermore, it is also equivalent to finding the eigenvector of \mathbf{A}^T associated to the eigenvalue $\lambda = 1$ and renormalizing it (i. e divide each element of the vector by the sum of all of its elements)

```
A = TransMat(adjMat = X, epsilon = 0.05)
v = eigen(t(A))$vectors[, 1]
Pi = Re(v)/sum(Re(v))
round(Pi, 3)
```

```
## [1] 0.138 0.140 0.097 0.185 0.095 0.138 0.128 0.078
```

Let's simulate a sequence of 1000 clicks of the random web surfer (we will start from webpage 2).

```
set.seed(20222023)
n = 1000
Sequence = rep(0, n)
Sequence[1] = 2
for (i in 2:n){
  Sequence[i] = sample(1:8, size = 1, prob = A[Sequence[i-1],])
}
```

- Let's estimate the transition matrix $\hat{\mathbf{A}}$.

```
temp = table(Sequence[1:n-1], Sequence[2:n])
A_hat = temp/rowSums(temp)
A_hat
```

```
##
##           1           2           3           4           5           6
##  1 0.007751938 0.418604651 0.000000000 0.279069767 0.000000000 0.000000000
##  2 0.206666667 0.026666667 0.280000000 0.000000000 0.253333333 0.233333333
##  3 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
##  4 0.000000000 0.327956989 0.349462366 0.322580645 0.000000000 0.000000000
##  5 0.000000000 0.344827586 0.000000000 0.000000000 0.011494253 0.000000000
##  6 0.208633094 0.000000000 0.000000000 0.287769784 0.194244604 0.309352518
##  7 0.318840580 0.000000000 0.000000000 0.369565217 0.000000000 0.311594203
##  8 0.380952381 0.000000000 0.000000000 0.000000000 0.333333333 0.285714286
##
##           7           8
```

```
## 1 0.000000000 0.294573643
## 2 0.000000000 0.000000000
## 3 1.000000000 0.000000000
## 4 0.000000000 0.000000000
## 5 0.356321839 0.287356322
## 6 0.000000000 0.000000000
## 7 0.000000000 0.000000000
## 8 0.000000000 0.000000000
```

- Let's estimate the stationary distribution and compare it to the one computed in the previous question. For comparison we will compute the \uparrow_1 norm of their difference.

```
v_hat = eigen(t(A_hat))$vectors[, 1]
Pi_hat = Re(v_hat)/sum(Re(v_hat))
round(Pi_hat, 3)
```

```
## [1] 0.129 0.149 0.107 0.188 0.087 0.139 0.138 0.063
```

```
print("The l1 norm of their difference is:")
```

```
## [1] "The l1 norm of their difference is:"
```

```
sum(abs(Pi - Pi_hat))
```

```
## [1] 0.06615326
```

```
print("The forward KL divergence is:")
```

```
## [1] "The forward KL divergence is:"
```

```
sum(Pi*log(Pi/Pi_hat))
```

```
## [1] 0.003661656
```

Web Communities

```
A_mat = matrix(c(0.7, 0.2, 0.1, 0.25, 0.7, 0.05, 0.1, 0.1, 0.8), byrow = T, nrow = 3)

B = matrix(c(0.2, 0, 0.1, 0, 0.1, 0.3, 0.1, 0.3, 0, 0.1, 0.2, 0.2, 0.2, 0, 0.1, 0.3, 0, 0, 0.1, 0.1, 0.1), byrow = T, nrow = 3)

hidden_states = c('Sport', 'Culture', 'Beauty')

emissions = c('Abdominaux', 'Cosmetiques', 'Livres', 'Age', 'Force', 'Endurance', 'Resilience', 'Creme')
```

```

new_B=matrix(nrow=55, ncol = 3)
B_couples=matrix(nrow=55, ncol = 3)

k=1
#all couples
for (i in 1:10){
  for (j in 1:10){
    if (i<=j){
      new_B[k, ]=B[i, ]*B[j, ]

      k=k+1

    }

  }
}

#Taking care of repeated couple (by brute force)
mask=c(1, 11, 20, 28, 35, 41, 46, 50, 53, 55)
B_couples[mask,]=new_B[mask,]
B_couples[2:10, ]=2*new_B[2:10, ]
B_couples[12:19, ]=2*new_B[12:19, ]
B_couples[21:27, ]=2*new_B[21:27, ]
B_couples[29:34, ]=2*new_B[29:34, ]
B_couples[36:40, ]=2*new_B[36:40, ]
B_couples[42:45, ]=2*new_B[42:45, ]
B_couples[47:49, ]=2*new_B[47:49, ]
B_couples[51:52, ]=2*new_B[51:52, ]
B_couples[54, ]=2*new_B[54, ]
colSums(B_couples)#it sums to 1

```

From the emission matrix of key words to that of couples

```
## [1] 1 1 1
```

Sampling from the model(HMM)

```

set.seed(20222023)
n_couples = 55
n_states = 3
Z = matrix(rep(0, 30*100), nrow = 30)
XX = matrix(rep(0, 30*100), nrow = 30)
pi_start = c(0.1, 0.3, 0.6)
for (i in 1:30){
  Z[i, 1] = sample(1:n_states, size = 1, prob = pi_start)
  XX[i, 1] = sample(1:n_couples, size = 1, prob = B_couples[, Z[i, 1]])
}

for (i in 1:30){
  for (j in 2:100){

```

```

Z[i, j] = sample(1:n_states, size = 1, prob = A_mat[Z[i, j-1], ])
XX[i, j] = sample(1:n_couples, size = 1, prob = B_couples[, Z[i, j]])
}
}

```

Let's implement the Baum-Welch algorithm and use it to estimate the parameters of the model using the data simulated from the previous question.

```

#implementation of the logsumexp trick
logsumexp <- function(logx) {
  # compute \log(\sum exp(logx)) by rescaling it by m = \max(logx)
  # indeed : \log(\sum exp(logx)) = m + \log(\sum exp(logx - m))
  # This ensures an exp(0) somewhere in the sum
  m = max(logx)
  return(m + log(sum(exp(logx - m))))
}

#The forward algorithm as explained in the appendix
forward = function(emis, trans, initial, data){
  #A function that implements the forward algorithm
  #initialization
  K = ncol(trans)
  T = ncol(data)
  n = nrow(data)
  epsilon = 1e-300#to avoid taking log of 0 if it happens
  logalph = array(rep(0, K*T*n), dim = c(n, K, T))

  for (i in 1:n){
    #fill in the first element in logspace
    logalph[i, , 1] = log(initial) + log(emis[data[i, 1], ] + epsilon)
  }

  for (t in 2:T){
    for (k in 1:K){
      for (i in 1:n){
        #the argument for the log sum exp function
        arg = log(trans[, k]+epsilon) + logalph[i, , t-1]
        logalph[i, k, t] = log(emis[data[i, t], k]+epsilon) + logsumexp(arg)
      }
    }
  }

  return(logalph)
}

#The backward algorithm as explained in the appendix
backward = function(emis, trans, initial, data){
  #A function that implements the backward algorithm
  #initialization
  K = ncol(trans)
  T = ncol(data)

```

```

n = nrow(data)
epsilon = 1e-300#to avoid taking log of 0 if it happens
logbeta = array(rep(0, K*T*n), dim = c(n, K, T))

for (i in 1:n){
  #the last element since beta_T=1 its log is 0
  logbeta[i, , T] = 0 + epsilon
}

for (t in (T-1):1){
  for (k in 1:K){
    for (i in 1:n){
      #the argument for the log sum exp function
      arg = log(trans[k, ]+epsilon) + log(emis[data[i, t+1], ]+epsilon) + logbeta[i, ,t+1]
      #fill in logbeta
      logbeta[i, k, t] = logsumexp(arg)
    }
  }
}
return(logbeta)
}

#Computing the \log\mathbf{x}_i matrix as explained in the appendix
logKSI = function(emis, trans, initial, data, fr, br){
  #this function computes the log ks is
  #initialization
  K = ncol(trans)
  T = ncol(data)
  n = nrow(data)
  epsilon = 1e-300
  KSI_mat = array(rep(0, n*K*K*(T-1)), dim = c(n, (T-1), K, K))
  #compute the loglikelihood of each sequence
  logP = apply(fr[, , T], 1, logsumexp)
  for (l in 1:K){
    for (k in 1:K){
      for (t in 1:(T-1)){
        for (i in 1:n){
          #fill in the ksi matrix
          KSI_mat[i, t, l, k] = (fr[i, l, t] + br[i, k, t+1] + log(trans[l, k]+epsilon) +
            log(emis[data[i, t+1], k]+epsilon) - logP[i])
        }
      }
    }
  }
  return(KSI_mat)
}

#all together we get the Baum-Welch algorithm
Baum_Welch = function(emis, trans, initial, data, n_it, epsilon=1e-6){

```

```

K = ncol(trans)
T = ncol(data)
n = nrow(data)
brow = nrow(emis)

# Pi_new = rep(0, K)
A_new = matrix(rep(0, K*K), nrow = K)
B_new = matrix(rep(0, brow*K), nrow = brow)
fr = forward(emis, trans, initial, data)
br = backward(emis, trans, initial, data)
KSI_mat = logKSI(emis, trans, initial, data, fr, br)

A_old = trans
B_old = emis
Pi_old = initial
loglik = c(sum(apply(fr[, , T], 1, logsumexp)))

comp = array(rep(0, n*K*T), dim = c(n, K, T))
comp[, , 1:(T-1)] = aperm(apply(KSI_mat, c(1, 2, 3), logsumexp), c(1, 3, 2))
comp[, , T] = apply(KSI_mat[, (T-1), , ], c(1, 3), logsumexp)
comp = exp(comp)
for (iter in 1:n_it){
  #a vectorized formula for pi
  Pi_new = exp(apply(fr[, , 1] + br[, , 1] - apply(fr[, , T], 1, logsumexp), 2, logsumexp)-log(n))
  denom = apply(exp(KSI_mat), c(1, 2, 3), sum)
  nume = exp(KSI_mat)

  # comp = exp(LogTau(B_old , A_old, Pi_old , data, fr, br))#for emission matrix

  for (l in 1:K){
    for(k in 1:K){
      A_new[l, k] = sum(nume[, , 1, k])/sum(denom[ , , 1])
    }
  }
  for (j in 1:brow){
    for (k in 1:K){
      B_new[j, k] = sum(comp[, k, ]*(XX==j))/sum(comp[, k, ])
    }
  }

  B_old = B_new
  A_old = A_new
  Pi_old = Pi_new
  fr = forward(B_old, A_old, Pi_old, data)
  br = backward(B_old, A_old, Pi_old, data)
  KSI_mat = logKSI(B_old, A_old, Pi_old, data, fr, br)
  loglik = c(loglik, sum(apply(fr[, , T], 1, logsumexp)))
  comp[, , 1:(T-1)] = aperm(apply(KSI_mat, c(1, 2, 3), logsumexp), c(1, 3, 2))
  comp[, , T] = apply(KSI_mat[, (T-1), , ], c(1, 3), logsumexp)
  comp = exp(comp)
  if (abs((loglik[iter+1] - loglik[iter])/loglik[iter])<= epsilon){
    break
  }
}

```



```

    }
}

param = list()
param$Pi = round(Pi_old, 2)
param$A = round(A_old, 2)
param$B = round(B_old, 2)
param$loglik = loglik
return(param)
}

```

Estimating the parameters using the Baum_Welch


```

#A naive Initialization
K = n_states
J = n_couples
init_pi = rep(1/K, K)
init_A = abs(matrix(rnorm(n=K^2, sd=10), K, K))
init_A = init_A/rowSums(init_A)
init_B = abs(matrix(rnorm(n=K*J, sd=10), K, J))
init_B = init_B/rowSums(init_B)
init_B = t(init_B)

#estimation

Output = Baum_Welch(emis = init_B, trans=init_A, initial = init_pi, data = XX, n_it=500)
plot(Output$loglik)

```




```
#An initialization by perturbing the initial parameters a little bit

#coded in class

perturbate_matrix = function(X, noise=7e-2) {
  # function to add small noise to the true parameter A & B
  X_pert = X +
    matrix(rnorm(prod(dim(X)), 0, sd=noise),
      nrow(X),
      ncol(X))
  X_pert = abs(X_pert) # enforce positive values !
  X_pert = X_pert / rowSums(X_pert)
  return(X_pert)
}

init_pi = (pi_start + abs(rnorm(n = 3, sd = 7e-2)))/sum((pi_start + abs(rnorm(n = 3, sd = 7e-2))))
init_A = perturbate_matrix(X = A_mat)
init_B = t(perturbate_matrix(X=t(B_couples)))
Output = Baum_Welch(emis = init_B, trans=init_A, initial = init_pi, data = XX, n_it=500)

plot(Output$loglik)
```



#we see that in this case it is better

Let's use Viterbi's algorithm to estimate the successive hidden states of the first row in the simulated data

```
#The Viterbi algorithm algorithm
param_test = list(A = A_mat, B = t(B_couples), pi = pi_start)
Viterbi<-function(x, param){
  epsilon<-1e-6
  K<-nrow(param$A)
  n<-length(x)
  S<-matrix(0,K,n)
  logV<-matrix(-Inf,K,n)
  Zest<-rep(0,n)
  for (k in 1:K){
    logV[k,1]<-log(param$B[k, x[1]]+epsilon)+log(param$pi[k])
  }
  # Forward
  for (t in (2:n))
    for (k in (1:K)){
      logV[k,t]=max(logV[,t-1]+log(param$A[,k])+log(param$B[k, x[t]]))
      S[k,t-1]=which.max(logV[,t-1]+log(param$A[,k])+log(param$B[k, x[t]]))
    }
  # Back-tracking
```

```

Zest[n]<-which.max(logV[,n])

for (t in (n-1):1)
Zest[t]<-S[Zest[t+1],t]
return(Zest)
}

Z1_hat = Viterbi(x=XX[1, ], param=param_test)

```

The percentage of errors made in the estimation of Z1 is

```
100*(sum(Z1_hat != Z[1, ])/length(Z1_hat))
```

```
## [1] 6
```

We observe above that we only made 6 errors, out the 100 estimated values. Which is not bad at all (using the optimal values).

Community Detection

Let there be 90 pages clustered in three groups $\{S, C, B\}$ we will note the $\{1, 2, 3\}$ Let's simulate an adjacency matrix of the directed graph using the stochastic block model described in the report

```

set.seed(20222023)
SBM_sample = function (n=90, pi_start = c(1/3, 1/3, 1/3), alpha=0.15, beta=0.05, loop=FALSE, directed=T)

  K = length(pi_start)

  Z = sample(1:K, replace = T, size = n, prob = pi_start)
  X = matrix(rep(0, n*n), nrow = n)
  Gamma = matrix(rep(alpha, K*K), nrow=K)
  Gamma[upper.tri(Gamma)|lower.tri(Gamma)] = beta

  #if the graph is directed
  if (directed) {

    #if loops are allowed (nodes can point to themselves)
    if (loop) {
      for (i in 1:n){
        for (j in 1:n){
          X[i, j] = rbinom(1, 1, prob = Gamma[Z[i], Z[j]])
        }
      }
    }

    #if loops are not allowed (nodes cannot point to themselves)
    else{
      for (i in 1:n){
        for (j in 1:n){
          if (i!=j) { X[i, j] = rbinom(1, 1, prob = Gamma[Z[i], Z[j]]) }
        }
      }
    }
  }

```

```

    }
  }
}

#if the graph is not directed(thus in this case X is symetric, and loop are not allowed)
else {
  for (i in 1:n){
    for (j in 1:n){
      if (i<j){
        X[i, j] = rbinom(1, 1, prob = Gamma[Z[i], Z[j]])
        X[j, i] = X[i, j]
      }
    }
  }
}

simul = list()
simul$Adjacency = X
simul$cluster = Z
return (simul)

}

gra = SBM_sample()

```

Plot the graph

```

#pdf('sbm_comm.pdf')
plot(graph_from_adjacency_matrix(gra$Adjacency,mode="directed"),vertex.color=gra$cluster)

```

```
#dev.off()
```

Let's compute the A^1 , and A^2

```
eps = 1/1000
A_one = gra$Adjacency + eps
A_one = A_one / rowSums(A_one)

A_two = (1/90)*matrix(rep(1, 90*90), nrow=90)
rowSums(A_one)
```

[illegible]

```
rowSums(A_two)
```

[illegible]

```
###Let's sample a sequence of 500 sequence of couples from the model with A_1
```

```

set.seed(20222023)
# websites are labeled from 1 to 90
# couples are labeled from 1 to 55 this is to facilitate the computations as explained in the report
group = gra$cluster
Web_1 = rep(0, 500)#webs using A1
cpls_1 = rep(0, 500)#couple using A1
Web_1[1]=7#we begin at website 7
cpls_1[1]=sample(1:n_couples, size=1, prob = B_couples[, group[Web_1[1]]])
for (i in 2:500){
  Web_1[i]=sample(1:90, size=1, prob = A_one[Web_1[i-1], ])
  cpls_1[i]=sample(1:55, size=1, prob = B_couples[, group[Web_1[i]]])
}

```

###Let's sample a sequence of 500 sequence of couples from the model with A_2

```

set.seed(20222023)
# websites are labeled from 1 to 90
# couples are labeled from 1 to 55 this is to facilitate the computations as explained in the report
group = gra$cluster
Web_2 = rep(0, 500)#webs using A2
cpls_2 = rep(0, 500)#couple using A2
Web_2[1]=7#we begin at website 7
cpls_2[1]=sample(1:n_couples, size=1, prob = B_couples[, group[Web_2[1]]])
for (i in 2:500){
  Web_2[i]=sample(1:90, size=1, prob = A_two[Web_2[i-1], ])
  cpls_2[i]=sample(1:55, size=1, prob = B_couples[, group[Web_2[i]]])
}

```